

Learning a function on a manifold

December 2025

1 Definitions and Previous Results

Definition 1. A metric tensor on (the embedded in $\mathbb{R}^k$) manifold $\mathcal{M}$ is a function $g : \mathcal{M} \rightarrow \mathbb{R}^{k \times k}, \forall x \in \mathcal{M}, \text{rank } g(x) = \dim \mathcal{M}, g(x) \succeq 0$.

Definition 2 (Dirichlet energy). The Dirichlet energy of a function $f : \mathcal{M} \rightarrow \mathbb{R}$ with respect to the manifold $\mathcal{M}$ equipped with the metric tensor g is defined as

$$\mathcal{E}_{\mathcal{M},g}(f) = \frac{1}{2} \int_{\mathcal{M}} \|\nabla_g f(x)\|_g^2 d\mu_g(x),$$

where $\nabla_g f$ denotes the gradient of f with respect to the metric g , $\|\cdot\|_g$ denotes the norm induced by g , and μ_g is the volume measure associated to g on $\mathcal{M}$.

Definition 3 (Measure). A Riemannian metric on a d -dimensional manifold is a smoothly varying inner product

$$g_x : T_x M \times T_x M \rightarrow \mathbb{R}.$$

In local coordinates, it is represented by a positive-definite matrix $g(x) = [g_{ij}(x)]$. The metric induces the Riemannian volume measure, often denoted:

$$d\mu_g(x) = \sqrt{\det g(x)} dx^1 \cdots dx^d.$$

This is the analogue of “area = $\sqrt{\det g}$ ” in differential geometry. $d\mu_g(x) = \sqrt{\det g(x)} dx^1 \cdots dx^d$.

Definition 4 (Gradient). The gradient of a function $f : \mathcal{M} \rightarrow \mathbb{R}$ with respect to the metric tensor g is defined as

$$\nabla_g f(x) = g(x) \nabla f(x) \in T_x \mathcal{M},$$

where $\nabla f(x)$ is the gradient of f with respect to the Euclidean metric.

Definition 5 (Laplacian). The Laplacian of a function $f : \mathcal{M} \rightarrow \mathbb{R}$ with respect to the metric g is defined as

$$\Delta_g f(x) = \nabla \cdot \nabla_g f(x) = \nabla \cdot (g \nabla f(x)) \in \mathbb{R}$$

Theorem 1 (Strong Whitney embedding theorem). The strong Whitney embedding theorem states that any smooth real m -dimensional manifold (required also to be Hausdorff and second-countable) can be smoothly embedded in the real $2m$ -space, $\mathbb{R}^{2m}$, if $m > 0$. This is the best linear bound on the smallest-dimensional Euclidean space that all m -dimensional manifolds embed in, as the real projective spaces of dimension m cannot be embedded into $(2m - 1)$ -space if m is a power of two (as can be seen from a characteristic class argument, also due to Whitney).

Theorem 2 (weak Whitney embedding theorem). The weak Whitney embedding theorem states that any continuous function from an n -dimensional manifold to an m dimensional manifold may be approximated by a smooth embedding provided $m > 2n$. Whitney similarly proved that such a map could be approximated by an immersion provided $m > 2n - 1$. This last result is sometimes called the Whitney immersion theorem.

2 Learning a function

Given:

- An embedded manifold $\mathcal{M} \subset \mathbb{R}^k, \dim \mathcal{M} = d$
- A metric tensor on $\mathcal{M}$: $g : \mathcal{M} \rightarrow \mathbb{R}^{k \times k}, \forall \mathbf{x} \in \mathcal{M}, \text{rank } g(\mathbf{x}) = d, g(\mathbf{x}) \succeq 0$.
Note: For simplicity, we assume $\det g(\mathbf{x}) = 1$ hence, the measure is the Lebesgue measure.
- A dataset $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ where $\mathbf{x}_i \in \mathcal{M}, y_i \in \mathbb{R}$

Goal: to learn the smoothest function $f : \mathcal{M} \rightarrow \mathbb{R}$ such that the function $f(\mathbf{x}_i) \approx y_i$.

Constrained optimization:

$$\begin{aligned} & \min_f \mathcal{E}_{\mathcal{M},g}(f) \\ \text{s.t. } & f(\mathbf{x}_i) = y_i, \forall i \in [n]. \end{aligned}$$

Unconstrained optimization:

$$\begin{aligned} & \min_f \mathcal{E}_{\mathcal{M},g}(f) + \lambda \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2 = \\ & \min_f \frac{1}{2} \int_{\mathcal{M}} \|\nabla_g f(\mathbf{x})\|_g^2 d\mu_g(\mathbf{x}) + \lambda \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2 \end{aligned}$$

Gradient descent, convex optimization (in a *function* space):

$$\begin{aligned} \text{PDE: } & \frac{\partial f}{\partial t} = -\Delta_g f + 2\lambda \sum_{i=1}^n (f(\mathbf{x}_i) - y_i) \delta_{\mathbf{x}_i} \\ \text{step: } & f_{t+1} = f_t + \eta \left(\Delta_g f_t - 2\lambda \sum_{i=1}^n (f(\mathbf{x}_i) - y_i) \delta_{\mathbf{x}_i} \right) \end{aligned}$$

Definition 6 (Gradient-flow evolution operator). We will denote $\Phi_T(f_i)$ as the gradient-flow evolution operator of f_i after time T of gradient descent. Note, that $\Phi_T(f_i) = \Phi_T(f_i, g, \mathcal{M}, D)$ depends on the metric tensor g , the manifold $\mathcal{M}$, and the dataset $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$.

Discretization in space (Disclaimer: this is a bit of a stretch.) If $\mathcal{M}$ is discretized as a point cloud $\mathcal{P} = \{\mathbf{z}_i\}_{i=1}^n$ with weights

$$w_{ij} = \exp \left(-\frac{\|\mathbf{z}_i - \mathbf{z}_j\|_{g(\mathbf{z}_i)}^2}{\epsilon} \right)$$

(epsilon TBD), one may argue that:

$$L_{\mathcal{P}} D_{\mathcal{P}}^{-1} \approx \Delta_g$$

Hence, one step of gradient descent becomes (yes, here comes the "stretch") like a transformer layer:

$$f_{t+1}(\mathbf{z}) = \underbrace{f_t(\mathbf{z})}_{\text{Residual}} + \eta \left(\underbrace{\frac{\sum_i \exp(-\|\mathbf{z} - \mathbf{z}_i\|_{g(\mathbf{z})}^2/\epsilon) f_t(\mathbf{z}_i)}{\sum_j \exp(-\|\mathbf{z} - \mathbf{z}_j\|_{g(\mathbf{z})}^2/\epsilon)}}_{\approx \text{Attention}} - \underbrace{2\lambda \sum_{i=1}^n (f(\mathbf{x}_i) - y_i) \delta_{\mathbf{x}_i}(\mathbf{z})}_{\text{Residual}} \right)$$

Idea 1: what if a fixed transformer performs gradient descent on the function space, hence, learns a function in context?

3 Learning a manifold

What if we are given a few smooth functions $f_1, \dots, f_l : \mathbb{R}^k \rightarrow \mathbb{R}$ on a manifold $\mathcal{M}$. However, we don't know the metric tensor g on $\mathcal{M}$. Can we learn the metric tensor such that the functions $f_1, \dots, f_l$ are as smooth as possible? If we do so, then we can easily learn a new smooth function from a few observations on the fly next time.

For this problem, the objective is the following:

$$\begin{aligned} & \min_g \sum_{i=1}^l \mathcal{E}_{\mathcal{M},g}(f_i) \\ \text{s.t. } & \det g(\mathbf{x}) = 1 \text{ for all } \mathbf{x} \in \mathcal{M}. \end{aligned}$$

where $\mathcal{E}_{\mathcal{M},g}(f_i)$ is the Dirichlet energy of f_i with respect to the metric tensor g . If we discard the constraint on the determinant or relax it to $\alpha I \preceq g(\mathbf{x}) \preceq \beta I$, the problem is convex w.r.t. g .

Okay, as of now, we didn't have a need to perform GD in a function space inside the global optimization with respect to the metric tensor g – there was no need to train a transformer-like model with multiple layers. In what case would that be necessary?

4 Lessons from Kernel Methods

- If we know *the right* kernel $\kappa(x, y)$, given $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ we can approximate any $f(x)$ in one step. For example, in Jacot, Gabriel, and Hongler [JGH20] we learn a function that would give us the right kernel.
- The generalization error will on
 - the number of samples n
 - the spectrum of $\mathcal{T}_\kappa$
 - the regularity of the f
 as the risk is expressed as

$$\mathcal{E}(\hat{f}_\lambda) = \underbrace{\sum_{j \geq 1} \left(\frac{\lambda}{\mu_j + \lambda} \right)^2 f_j^2}_{\text{bias term}} + \underbrace{\frac{\sigma^2}{n} \sum_{j \geq 1} \frac{\mu_j}{(\mu_j + \lambda)^2}}_{\text{variance term}}$$

- So again, given some functions $f_1, \dots, f_i$ we want to find kernel κ with respect to what those functions are as smooth as possible.
- If we can do that we can easily learn a function on a fly.
- paper Shen et al. [She+25] bridges kernel methods to $O(1)$ -layer transformers.

5 Distorted Manifold Hypothesis

Let μ be a probability distribution supported on a d -dimensional manifold $\mathcal{M} \subset \mathbb{R}^k$.

Hypothesis: True data: set of points $x \in \mathbb{R}^k$ that comes from μ – a probability distribution supported on a d -dimensional manifold $\mathcal{M} \subset \mathbb{R}^k$, $d \ll k$. However, we some distortion process $Z(\mathbf{x}, t)$ we observed distorted points $\tilde{\mathbf{x}}$ not $\mathbf{x}$. For example, the distortion governed by time-inhomogeneous diffusion $D(t) : [0, 1] \rightarrow \mathbb{R}^{n \times n}$, $D \succeq 0$ and drift $b : [0, 1] \rightarrow \mathbb{R}^n$.

$$\begin{cases} \frac{\partial \mu_t(x)}{\partial t} = \nabla_x \cdot (D(t) \nabla_x \mu_t(x)) - b(t) \cdot \nabla_x \mu_t(x), & x \in \mathbb{R}^k, t \in (0, 1], \\ \mu_{t=0}(x) = \mu(x), & x \in \mathbb{R}^k. \end{cases}$$

For an effective kernel method, we should eliminate some distortion before fitting observed data: $\{(\tilde{\mathbf{x}}_i, y_i)\}$. I believe that finding *minimal-embedding* of μ helps with that.

Definition 7. A *minimal embedding* of μ is the map

$$\Phi : \mathcal{M} \rightarrow \mathbb{R}^k$$

that minimizes the Dirichlet energy of the embedding:

$$\Phi^* = \arg \min_{\Phi : \mathcal{M} \rightarrow \mathbb{R}^k} \mathcal{E}(\Phi) := \frac{1}{2} \sum_{j=1}^k \int_{\mathcal{M}} \|\nabla_{\mathcal{M}} \Phi_j(x)\|^2 d\mu(x)$$

subject to natural constraints:

- Φ preserves the intrinsic dimension (rank of the Jacobian is d almost everywhere),
- $\mathbb{E}_\mu[\Phi(x)] = 0$ (remove trivial translations),
- $\frac{1}{d} \int_{\mathcal{M}} \Phi(x) \Phi(x)^\top d\mu(x) = I_d$

Intuition: the minimal embedding is the least energetic or smoothest way to embed the distribution into $\mathbb{R}^k$. Equivalently:

- it is the embedding whose coordinate functions are harmonic: $\Delta_{\mathcal{M}} f = 0$.
- it removes all geometric distortions caused by drift/diffusion processes,
- it recovers the "canonical" coordinates of the underlying manifold w.r.t. the intrinsic geometry induced by μ .
- Minimal embedding = the embedding with the simplest geometry for which μ looks undistorted.

Idea: Transformer layers reverse the distortion and recovered the minimal energy of the $\tilde{\mu}$.

For example to reverse the diffusion one should do:

$$\begin{cases} -\partial_t \mu_t(x) = -\nabla \cdot (b_t(x) \mu_t(x)) + \nabla \cdot (D_t(x) \nabla \mu_t(x)), & t \in [0, 1], \\ \mu_1(x) = \tilde{\mu}(x). \end{cases}$$

One transformer layer = One gradient step in the measure space that captures the equation above (on a discretized μ).

However, this process is ill-posed as the diffusion destroys high-frequency models so some information will be inevitably lost.

Relevant literature

- Geshkovski et al. [Ges+25] – models transformers as interacting particle systems
- Belkin and Niyogi [BN06] – on the convergence of discrete Laplacian to Laplace-Beltrami operator
- Zhang and Zhou [ZZ25] – PDE view on transformers
- Jacot, Gabriel, and Hongler [JGH20]

References

- [BN06] Mikhail Belkin and Partha Niyogi. “Convergence of Laplacian Eigenmaps”. In: *Advances in Neural Information Processing Systems*. Vol. 19. MIT Press, 2006. (Visited on 12/05/2025).
- [Ges+25] Borjan Geshkovski et al. *A Mathematical Perspective on Transformers*. Aug. 2025. DOI: [10.48550/arXiv.2312.10794](https://doi.org/10.48550/arXiv.2312.10794). arXiv: [2312.10794](https://arxiv.org/abs/2312.10794) [cs]. (Visited on 12/05/2025).
- [JGH20] Arthur Jacot, Franck Gabriel, and Clément Hongler. *Neural Tangent Kernel: Convergence and Generalization in Neural Networks*. Feb. 2020. DOI: [10.48550/arXiv.1806.07572](https://doi.org/10.48550/arXiv.1806.07572). arXiv: [1806.07572](https://arxiv.org/abs/1806.07572) [cs]. (Visited on 08/21/2025).
- [She+25] Zhaiming Shen et al. *Understanding In-Context Learning on Structured Manifolds: Bridging Attention to Kernel Methods*. Oct. 2025. DOI: [10.48550/arXiv.2506.10959](https://doi.org/10.48550/arXiv.2506.10959). arXiv: [2506.10959](https://arxiv.org/abs/2506.10959) [cs]. (Visited on 12/08/2025).
- [ZZ25] Yukun Zhang and Xueqing Zhou. *PDE-Transformer: A Continuous Dynamical Systems Approach to Sequence Modeling*. Oct. 2025. DOI: [10.48550/arXiv.2510.03272](https://doi.org/10.48550/arXiv.2510.03272). arXiv: [2510.03272](https://arxiv.org/abs/2510.03272) [cs]. (Visited on 12/05/2025).