An Energy-Based View on the Manifold Hypothesis

Anonymous Author(s)
Affiliation
Address
`email`

Abstract

A commonly adopted assumption in modern machine learning is the *manifold hypothesis*, which claims that although data is embedded in a high-dimensional ambient space, the probability mass concentrates near a much lower-dimensional geometric structure. However, the real-world latent variables that images are drawn from are continuous, not discrete, and therefore cannot be completely captured by a finite manifold. Moreover, the assumption does not properly capture the union of manifolds, stratified spaces, and manifolds with singularities and boundaries. Drawing upon recent advances in energy-based models, we provide an energy-based perspective on the real-world latents, describing an Occam's razor style analog for the manifold hypothesis, which is more amenable to recent literature in the field. We theoretically prove that our construction recovers the manifold hypothesis for coarse functions and provide a diffusion-based perspective on it. Finally, we validate our theoretical framework with numerical examples.

1 Introduction

If pictures are discretized objects in $\mathbb{R}^d$, where $d = n^2$ is the dimension of the image.

We consider raw objects $f : [0,1]^2 \to \mathbb{R}^3$ corresponding to depixelated RGB pictures.

A commonly adopted assumption in modern machine learning is the *manifold hypothesis*, which claims that although data is embedded in a high-dimensional ambient space, the probability mass concentrates near a much lower-dimensional geometric structure.

Assumption 1.1 (Manifold Hypothesis)**.** Let $\chi \subset \mathbb{R}^D$ be the ambient space and let $\mathcal{D}$ be a probability distribution on χ. There exists a d-dimensional manifold $\mathcal{M} \subset \chi$, with $d \ll D$, such that

$$\mathbb{P}_{x\sim\mathcal{D}}\big(\operatorname{dist}(x,\mathcal{M}) \leq \varepsilon\big) \geq 1-\delta,$$

for sufficiently small $\varepsilon, \delta > 0$.

In practice, data rarely concentrate on a single smooth manifold. Rather, they are better modeled as lying on *unions of manifolds*, or more generally on *stratified spaces* composed of pieces of varying intrinsic dimensions, potentially with singularities and boundaries.

Why small d is important?

- Generalization error depends on the intrinsic dimension d of the manifold.
- The number of samples n needed to learn the manifold increases with d.
- The complexity of the model increases with d.
- Can get a small embedding.

Some shortcomings of the manifold hypothesis are that

- Doesn't properly capture union of manifolds.
- Doesn't properly capture stratified spaces.
- Doesn't properly capture manifolds with singularities and boundaries.
- What about infinite-dimensional data? Cause images come from infinite-dimensional continuous functions.

[Emile says: would be nice to elaborate on the above, and to point to specific instances where assuming it has some "catastrophic" consequences. that could go into a motivating examples section]

Data may be infinite-dimensional (images as functions), unions/strata, etc., so a "single smooth manifold" is an inadequate representation.

Contributions:

1. We propose an effective-dimension hypothesis at finite scale, and connect it to Occam penalties in energy-based learning,
2. We define a multiscale effective dimension hypothesis on general Dirichlet spaces and finite discretizations (showing that coarse/discretized functions recover the manifold hypothesis),
3. We show that under quadratic/Laplacian regularization, an Occam term (log det Hessian) induces a trace/effective-dimension quantity that is scale-dependent, and it is large for noise/fractal-like signals,
4. We show something diffusion related

2 Preliminaries

We work in an abstract setting that captures the common statistical belief that observed data $x_1, \ldots, x_n \in X$ concentrate on a low-dimensional structure, even when the ambient space X is high-dimensional. Concretely, we fix a *reference* distribution ν_0 on X (e.g. Gaussian or uniform) encoding our prior knowledge about X, and posit that the data are drawn from a *target* distribution μ that is far more concentrated — e.g. supported on a d-dimensional manifold $M \subset \mathbb{R}^D$ with $d \ll D$. The benefit is quantitative: greater concentration of μ yields faster learning rates for functions $f : X \to \mathbb{R}$.

Differential structure. To differentiate functions on X we fix an *algebra of test functions* $\mathcal{A}$ (a subset of $\mathbb{R}^X$ closed under pointwise products) and a *derivation*

$$D : \mathcal{A} \longrightarrow \mathcal{M}$$

into a module $\mathcal{M}$ equipped with a pointwise inner product $\langle \cdot, \cdot \rangle_x$, satisfying the Leibniz rule $D(fg) = fDg + gDf$. This gives rise to the *carré du champ* operator

$$\Gamma(f, g)(x) := \langle Df(x),\, Dg(x) \rangle_x,$$

which generalizes the square-gradient $|\nabla f|^2$ to the abstract setting.

Dirichlet form and Laplacian. Given any σ-finite measure ν on $(X, \mathcal{F})$, the carré du champ integrates to a *Dirichlet form*:

Definition 2.1 (Symmetric Dirichlet Space [Bakry et al., 2014, §3.4])**.** Let $(X, \mathcal{F}, \mu)$ be a σ-finite measure space. A *Dirichlet form* on $L^2(\mu)$ is a pair $(\mathcal{E}, \mathcal{D}(\mathcal{E}))$ where:

1. $\mathcal{D}(\mathcal{E}) \subset L^2(\mu)$ is a dense linear subspace;
2. $\mathcal{E} : \mathcal{D}(\mathcal{E}) \times \mathcal{D}(\mathcal{E}) \to \mathbb{R}$ is a symmetric bilinear form;
3. **(Non-negativity)** $\mathcal{E}(f, f) \geq 0$ for all $f \in \mathcal{D}(\mathcal{E})$;
4. **(Closedness)** $\mathcal{D}(\mathcal{E})$ is complete under $\|f\|_{\mathcal{E}}^2 := \|f\|_{L^2(\mu)}^2 + \mathcal{E}(f, f)$;
5. **(Markov property)** for every $f \in \mathcal{D}(\mathcal{E})$, the truncation $\widetilde{f} := (0 \vee f) \wedge 1$ satisfies $\widetilde{f} \in \mathcal{D}(\mathcal{E})$ and $\mathcal{E}(\widetilde{f}, \widetilde{f}) \leq \mathcal{E}(f, f)$.

The quadruple $(X, \mathcal{F}, \mu, \mathcal{E})$ is called a *Symmetric Dirichlet Space*.

In our setting the Dirichlet form is

$$\mathcal{E}_\nu(f, g) = \int_X \Gamma(f, g)\, d\nu,$$

so $\mathcal{E}_\nu(f, f)$ is the total gradient energy of f under ν. The Markov property is automatic from the Leibniz rule and the truncation behaviour of Γ.

Definition 2.2 (Generalized Laplacian)**.** To any Dirichlet form $(\mathcal{E}_\nu, \mathcal{D}(\mathcal{E}_\nu))$ on $L^2(\nu)$ there corresponds a unique self-adjoint operator L_ν on $L^2(\nu)$ satisfying

$$\mathcal{E}_\nu(f, g) = \langle f,\, L_\nu g\rangle_{L^2(\nu)}, \qquad \forall\, f \in \mathcal{D}(\mathcal{E}_\nu),\ g \in \mathcal{D}(L_\nu).$$

L_ν admits a spectral decomposition $L_\nu \phi_j = \lambda_j \phi_j$ with

$$0 = \lambda_0 \le \lambda_1 \le \lambda_2 \le \cdots, \qquad \langle \phi_j, \phi_k\rangle_{L^2(\nu)} = \delta_{jk}.$$

The Markov property ensures that e^{-tL_ν} is a Markov semigroup: starting from any initial distribution ν_0 on X, the heat flow $e^{-tL_\nu}\nu_0$ converges to ν as $t \to \infty$, confirming that ν is the invariant measure of the diffusion. The *iterated carré du champ* Γ_2^ν measures how Γ evolves under this semigroup:

$$\Gamma_2^\nu(f, g) = \tfrac{1}{2}\left(L_\nu \Gamma(f, g) - \Gamma(f, L_\nu g) - \Gamma(g, L_\nu f)\right), \qquad \Gamma_2^\nu(f) := \Gamma_2^\nu(f, f).$$

The Bakry–Émery condition $\Gamma_2^\nu(f) \ge \rho\, \Gamma(f)$ ($\rho \in \mathbb{R}$) encodes a lower bound on the Ricci curvature of the underlying space.

Effective dimension. The eigenvalue asymptotics of L_ν encode the intrinsic dimension of the learning space, providing the key rate parameter throughout the paper.

Definition 2.3 (Effective Dimension)**.** The *effective dimension* of $(X, \mathcal{F}, \nu, \mathcal{E}_\nu)$ is the quantity $d_{\text{eff}} \in (0, \infty)$ characterised by any of the following equivalent asymptotics:

$$\begin{aligned} \lambda_j &\asymp j^{2/d_{\text{eff}}} && \text{(eigenvalue growth)},\\ N(\lambda) := \#\{j : \lambda_j \le \lambda\} &\asymp \lambda^{d_{\text{eff}}/2} && \text{(counting function)},\\ \operatorname{tr} e^{-tL_\nu} &\sim t^{-d_{\text{eff}}/2} && \text{(heat-kernel trace, as } t \to 0^+). \end{aligned}$$

For a compact d-dimensional Riemannian manifold M, Weyl's law gives $d_{\text{eff}} = d$, confirming that the effective dimension recovers the geometric dimension in the classical setting.

3 Effective Dimension Hypothesis

Assumption 3.1 (Effective Dimension Hypothesis)**.** All the observed data $x_1, \ldots, x_n \in X$ come from a learning space $(X, \mathcal{A}, \mu)$ with small effective dimension even if the ambient space $(X, \mathcal{A}, \nu_0)$ has large effective dimension where ν_0 is the base distribution on X:

$$d_{\text{eff}}(\mu) \ll d_{\text{eff}}(\nu_0)$$

Definition 3.2 (Curvature–Dimension Condition, Ledoux [2000] (Definition 1.2))**.** For $\rho \in \mathbb{R}, N \in (0, \infty)$, L_ν satisfies $\mathrm{CD}(\rho, N)$ if for all $f \in \mathcal{A}$

$$\Gamma_2^\nu(f, f) \ge \rho \Gamma(f, f) + \frac{1}{N}\left(L_\nu f\right)^2$$

Intuitively: $\rho \sim$ curvature parameter, $N \sim$ dimension

Example 3.3. *Examples of generators L_ν that satisfy the CD condition:*

- *An n-dimensional complete Riemannian manifold (M, g) with Ricci curvature bounded below, or rather the Laplacian Δ on M, satisfies the inequality $\mathrm{CD}(R, n)$ with R the infimum of the Ricci tensor and n the geometric dimension.*
- *The classical Laplacian Δ on $\mathbb{R}^n$ satisfies $\mathrm{CD}(0, n)$,*
- *the Laplace-Beltrami operator on the unit sphere S^n is of curvature dimension $\mathrm{CD}(n-1, n)$ since curvature is constant and equal to $n-1$ in this case.*

3.1 Recovering the Manifold Hypothesis upon coarsening

Moreover, we show that the effective dimension hypothesis recovers the manifold hypothesis when the infinite-dimensional space is coarsened into pixels.

Theorem 3.4. *Recovering the Manifold Hypothesis Let $M \subset \mathbb{R}^D$ be a compact embedded C^2 submanifold with tubular radius r_0, with the nearest-point projection $\pi : T_{r_0}(M) \to M$. Fix $0 < \epsilon < r_0$. Let μ be a Borel probability measure supported on the tangent space $T_\epsilon(M) \coloneqq \{x \in \mathbb{R}^D : \mathrm{dist}(x, M)\}$, and define the manifold-supported measure $\overline{\mu} \coloneqq \pi_{\#}\mu$. Then, for every $p \geq 1$, we have $W_p(\mu, \overline{\mu}) \leq \epsilon$ and for every L-Lipschitz observable function $f : \mathbb{R}^D \to \mathbb{R}$, we have that*

$$\left| \int f \mathrm{d}\mu - \int f \mathrm{d}\overline{\mu} \right| \leq L\epsilon.$$

Corollary 3.5. *Let $F_\delta \coloneqq \{f : \mathrm{Lip}(f) \leq 1/\delta\}$. By Theorem B.1, we have that $\sup_{f \in F_\delta} | \int f \mathrm{d}\mu - \int f \mathrm{d}\overline{\mu}| \leq \epsilon/\delta$. Therefore, once $\delta \gg \epsilon$, the tube supported distribution ϵ and the manifold-supported distribution $\overline{\mu}$ are indistinguishable by coarse observations. Therefore, the manifold hypothesis is recovered upon coarsening the effective dimension hypothesis.*

Remark 3.6. The proof of Theorem B.1 is based on the metric space, not on any of its spectral properties. Moreover, the manifold structure is only used to guarantee the existence of the projection π, and the same argument would also follow for any measurable set $S \subset \mathbb{R}^D$ equipped with a measurable map $\pi : T_\epsilon(S) \to S$ such that $\|x - \pi(x)\| \leq \epsilon$, which allows the result to extend beyond smooth manifolds. Finally, the Lipschitz bound in Theorem B.1 is tight in general. For this, take $M = \{0\} \subset \mathbb{R}^D$, $\mu = \delta_\epsilon, \overline{\mu} = \delta_0$, and $f(x) = Lx$. Then, $| \int f \mathrm{d}\mu - \int f \mathrm{d}\overline{\mu}| = L\epsilon$.

3.2 References

https://djalil.chafai.net/blog/2023/01/12/log-sobolev-and-bakry-emery/

3.3 Somewhat Related Works

Energy-based models Song and Kingma [2021], Du and Mordatch [2019], LeCun et al. [2006], Zhu et al. [2024], Du et al. [2021, 2020]

Compositional generation Su et al. [2024], Liu et al. [2021, 2022, 2023], Du and Kaelbling [2024], Bagautdinov et al. [2018]

Occam's razor Elmoznino et al. [2025], Rasmussen and Ghahramani [2000]

Bakry-Emery

Dirichlet Form

Example: Vision Data Each 2d image (or 3d scene or a point cloud) can be modeled $f : \mathbb{R}^2 \to \mathbb{R}^3$ as a function from coordinates to RGB (or XYZ) values. That's how NeRFS work:

Neural Radiance Fields (NeRFs). Mildenhall et al. [2021], Pumarola et al. [2021], Xu et al. [2021], Wang et al. [2022], Rabby and Zhang [2024], Reiser et al. [2021], Müller et al. [2022], Tancik et al. [2021], Yu et al. [2021], Chen et al. [2022] [Emile says: this actually looks very relevant. would be worth doing a deeper dive into this] A *Neural Radiance Field* models a 3D scene as a continuous function

$$f_\theta : (\mathbf{x}, \mathbf{d}) \;\mapsto\; (\sigma, \mathbf{c}),$$

where

- $\mathbf{x} \in \mathbb{R}^3$ denotes a spatial location,
- $\mathbf{d} \in \mathbb{S}^2$ denotes a viewing direction,
- $\sigma \in \mathbb{R}_{\geq 0}$ is the volume density at $\mathbf{x}$,
- $\mathbf{c} \in \mathbb{R}^3$ is the emitted RGB color in direction $\mathbf{d}$.

The key idea is that the scene is represented neither as a mesh, nor as a voxel grid, nor as a point cloud, but as a *continuous radiance field* parameterized by a neural network.

In a similar way, each 2d image can be modeled as a function from coordinates to RGB values: $f : \mathbb{R}^2 \to \mathbb{R}^3$ and the embedding of this function would the weights of the corresponding ReLU network. The regular manifold hypothesis doesn't properly capture this structure that:

- In a union of manifolds
- Has some hierarchical structure.
- Has permutation-invariance with respect to weights in each layer.

Also, importantly, each image f is not a smooth function: its a combination of a "simple" in some sense partition of the image into smooth enough regions.

4 Experiments

A learning curve experiment: estimate a function class regularized by the energy, and show that lower finite-scale complexity predicts sample complexity better than ambient dimension.

A 2D synthetic suite one smooth manifold, a union of two manifolds with intersection, and a stratified example with a singular point. Show that asymptotic Weyl dimension is not the right discriminator, while $d_\mu(t)$ at finite t separates them.

5 Conclusion

References

Timur Bagautdinov, Chenglei Wu, Jason Saragih, Pascal Fua, and Yaser Sheikh. Modeling facial geometry using compositional VAEs. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3877–3886, 2018.

Dominique Bakry, Ivan Gentil, and Michel Ledoux. *Analysis and Geometry of Markov Diffusion Operators*, volume 348 of *Grundlehren Der Mathematischen Wissenschaften*. Springer International Publishing, Cham, 2014. ISBN 978-3-319-00226-2 978-3-319-00227-9. doi: 10.1007/978-3-319-00227-9.

Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. Tensorf: Tensorial radiance fields. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXII*, page 333–350, Berlin, Heidelberg, 2022. Springer-Verlag. ISBN 978-3-031-19823-6. doi: 10.1007/978-3-031-19824-3_20. URL `https://doi.org/10.1007/978-3-031-19824-3_20`.

Yilun Du and Leslie Kaelbling. Position: compositional generative modeling: a single model is not all you need. In *Proceedings of the 41st International Conference on Machine Learning*, ICML'24. JMLR.org, 2024.

Yilun Du and Igor Mordatch. Implicit generation and modeling with energy based models. In *Advances in Neural Information Processing Systems*, volume 32, 2019.

Yilun Du, Shuang Li, and Igor Mordatch. Compositional visual generation with energy based models. *Advances in Neural Information Processing Systems*, 33:6637–6647, 2020.

Yilun Du, Shuang Li, Yash Sharma, Joshua B. Tenenbaum, and Igor Mordatch. Unsupervised learning of compositional energy concepts. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021.

Eric Elmoznino, Tom Marty, Tejas Kasetty, Leo Gagnon, Sarthak Mittal, Mahan Fathi, Dhanya Sridhar, and Guillaume Lajoie. In-context learning and occam's razor, 2025. URL `https://arxiv.org/abs/2410.14086`.

Yann LeCun, Sumit Chopra, Raia Hadsell, M Ranzato, Fujie Huang, et al. A tutorial on energy-based learning. *Predicting structured data*, 1(0), 2006.

Michel Ledoux. The geometry of Markov diffusion generators. *Annales de la Faculté des sciences de Toulouse : Mathématiques*, 9(2):305–366, 2000. ISSN 2258-7519.

Nan Liu, Shuang Li, Yilun Du, Joshua B. Tenenbaum, and Antonio Torralba. Learning to compose visual relations. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021.

Nan Liu, Shuang Li, Yilun Du, Antonio Torralba, and Joshua B. Tenenbaum. Compositional visual generation with composable diffusion models. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XVII*, page 423–439, Berlin, Heidelberg, 2022. Springer-Verlag. ISBN 978-3-031-19789-5. doi: 10.1007/978-3-031-19790-1_26. URL `https://doi.org/10.1007/978-3-031-19790-1_26`.

Nan Liu, Yilun Du, Shuang Li, Joshua B. Tenenbaum, and Antonio Torralba. Unsupervised compositional concepts discovery with text-to-image generative models. In *ICCV*, pages 2085–2095, 2023.

Ben Mildenhall, Pratul P Srinivasan, Matthew Tancik, Jonathan T Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. *Communications of the ACM*, 65(1):99–106, 2021.

Thomas Müller, Alex Evans, Christoph Schied, and Alexander Keller. Instant neural graphics primitives with a multiresolution hash encoding. *ACM Trans. Graph.*, 41(4), July 2022. ISSN 0730-0301. doi: 10.1145/3528223.3530127. URL `https://doi.org/10.1145/3528223.3530127`.

Albert Pumarola, Enric Corona, Gerard Pons-Moll, and Francesc Moreno-Noguer. D-nerf: Neural radiance fields for dynamic scenes. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10318–10327, 2021.

AKM Shahariar Azad Rabby and Chengcui Zhang. Beyondpixels: A comprehensive review of the evolution of neural radiance fields, 2024. URL `https://arxiv.org/abs/2306.03000`.

Carl Rasmussen and Zoubin Ghahramani. Occam's razor. *Advances in neural information processing systems*, 13, 2000.

Christian Reiser, Songyou Peng, Yiyi Liao, and Andreas Geiger. Kilonerf: Speeding up neural radiance fields with thousands of tiny mlps. In *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 14315–14325, 2021. doi: 10.1109/ICCV48922.2021.01407.

Yang Song and Diederik P. Kingma. How to train your energy-based models, 2021.

Jocelin Su, Nan Liu, Yanbo Wang, Joshua B. Tenenbaum, and Yilun Du. Compositional image decomposition with diffusion models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 46823–46842. PMLR, 21–27 Jul 2024.

Matthew Tancik, Ben Mildenhall, Terrance Wang, Divi Schmidt, Pratul P. Srinivasan, Jonathan T. Barron, and Ren Ng. Learned initializations for optimizing coordinate-based neural representations. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2845–2854, 2021. doi: 10.1109/CVPR46437.2021.00287.

Zirui Wang, Shangzhe Wu, Weidi Xie, Min Chen, and Victor Adrian Prisacariu. Nerf–: Neural radiance fields without known camera parameters, 2022. URL `https://arxiv.org/abs/2102.07064`.

Hongyi Xu, Thiemo Alldieck, and Cristian Sminchisescu. H-nerf: Neural radiance fields for rendering and temporal reconstruction of humans in motion. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021. URL `https://arxiv.org/pdf/2110.13746.pdf`.

Alex Yu, Vickie Ye, Matthew Tancik, and Angjoo Kanazawa. pixelnerf: Neural radiance fields from one or few images. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4576–4585, 2021. doi: 10.1109/CVPR46437.2021.00455.

Yaxuan Zhu, Jianwen Xie, Ying Nian Wu, and Ruiqi Gao. Learning energy-based models by cooperative diffusion recovery likelihood. In *The Twelfth International Conference on Learning Representations*, 2024.

A Mathematical Background

B Recovering the Manifold Hypothesis

We show that the manifold hypothesis is recovered after coarsening the effective dimension hypothesis in Theorem B.1.

Theorem B.1. *Recovering the Manifold Hypothesis Let $M \subset \mathbb{R}^D$ be a compact embedded C^2 submanifold with tubular radius r_0, with the nearest-point projection $\pi : T_{r_0}(M) \to M$. Fix $0 < \epsilon < r_0$. Let μ be a Borel probability measure supported on the tangent space $T_\epsilon(M) := \{x \in \mathbb{R}^D : \mathrm{dist}(x, M)\}$, and define the manifold-supported measure $\overline{\mu} := \pi_{\#}\mu$. Then, for every $p \geq 1$, we have $W_p(\mu, \overline{\mu}) \leq \epsilon$ and for every L-Lipschitz observable function $f : \mathbb{R}^D \to \mathbb{R}$, we have that*

$$\left| \int f \mathrm{d}\mu - \int f \mathrm{d}\overline{\mu} \right| \leq L\epsilon.$$

Proof. Let $X \sim \mu$ and set $Y := \pi(X)$. Since $\overline{\mu} := \pi_{\#}\mu$, we have $Y \sim \overline{\mu}$. Since $X \in T_\epsilon(M)$ almost surely, and $\pi(X)$ is the nearest point on M, we have $\|X - Y\| = \|X - \pi(X)\| = \mathrm{dist}(X, M) \leq \epsilon$, almost surely. Next, we consider the coupling $\gamma = (\mathrm{Id}, \pi)_{\#\mu} \in \Pi(\mu, \overline{\mu})$ of μ and $\overline{\mu}$, and γ is supported on pairs x, y with $\|x - y\| \leq \epsilon$. Therefore, for $1 \leq p < \infty$, we have

$$\begin{aligned} W_p(\mu, \overline{\mu})^p &\leq \int \|x - y\|^p \mathrm{d}\gamma(x, y) \\ &= \int \|x - \pi(x)\|^p \mathrm{d}\mu(x) \leq \epsilon^p, \end{aligned}$$

which implies $W_p(\mu, \overline{\mu}) \leq \epsilon$. Moreover, we have

$$W_\infty(\mu, \overline{\mu}) \leq \operatorname*{ess\,sup}_{x \sim \mu} \|x - \pi(x)\| \leq \epsilon.$$

Next, since $\overline{\mu} = \pi_{\#}\mu$, we have $\int f \mathrm{d}\overline{\mu} = \int f(\pi(x)) \mathrm{d}\mu(x)$. Therefore, we have

$$\begin{aligned} \left| \int f \mathrm{d}\mu - \int f \mathrm{d}\overline{\mu} \right| &= \left| \int (f(x) - f(\pi(x)) \mathrm{d}\mu(x) \right| \\ &\leq \int \|f(x) - f(\pi(x))\| \mathrm{d}\mu(x) \\ &\leq \int L\|x - \pi(x)\| \mathrm{d}\mu(x) \\ &\leq L\epsilon \int \mathrm{d}\mu(x) := L\epsilon, \end{aligned}$$

which proves the theorem. □