
An Energy-Based View on the Manifold Hypothesis

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A commonly adopted assumption in modern machine learning is the *manifold*
2 *hypothesis*, which claims that although data is embedded in a high-dimensional
3 ambient space, the probability mass concentrates near a much lower-dimensional
4 geometric structure. However, the real-world latent variables that images are drawn
5 from are continuous, not discrete, and therefore cannot be completely captured by
6 a finite manifold. Moreover, the assumption does not properly capture the union
7 of manifolds, stratified spaces, and manifolds with singularities and boundaries.
8 Drawing upon recent advances in energy-based models, we provide an energy-
9 based perspective on the real-world latents, describing an Occam’s razor style
10 analog for the manifold hypothesis, which is more amenable to recent literature
11 in the field. We theoretically prove that our construction recovers the manifold
12 hypothesis for coarse functions and provide a diffusion-based perspective on it.
13 Finally, we validate our theoretical framework with numerical examples.

14 1 Introduction

15 If pictures are discretized objects in $\mathbb{R}^d$, where $d = n^2$ is the dimension of the image.

16 We consider raw objects $f : [0, 1]^2 \rightarrow \mathbb{R}^3$ corresponding to depixelated RGB pictures.

17 A commonly adopted assumption in modern machine learning is the *manifold hypothesis*, which
18 claims that although data is embedded in a high-dimensional ambient space, the probability mass
19 concentrates near a much lower-dimensional geometric structure.

20 **Assumption 1.1** (Manifold Hypothesis). Let $\chi \subset \mathbb{R}^D$ be the ambient space and let $\mathcal{D}$ be a probability
21 distribution on χ . There exists a d -dimensional manifold $\mathcal{M} \subset \chi$, with $d \ll D$, such that

$$\mathbb{P}_{x \sim \mathcal{D}}(\text{dist}(x, \mathcal{M}) \leq \varepsilon) \geq 1 - \delta,$$

22 for sufficiently small $\varepsilon, \delta > 0$.

23 In practice, data rarely concentrate on a single smooth manifold. Rather, they are better modeled as
24 lying on *unions of manifolds*, or more generally on *stratified spaces* composed of pieces of varying
25 intrinsic dimensions, potentially with singularities and boundaries.

26 Why small d is important?

- 27 • Generalization error depends on the intrinsic dimension d of the manifold.
- 28 • The number of samples n needed to learn the manifold increases with d .
- 29 • The complexity of the model increases with d .
- 30 • Can get a small embedding.

31 Some shortcomings of the manifold hypothesis are that

- Doesn't properly capture union of manifolds.
- Doesn't properly capture stratified spaces.
- Doesn't properly capture manifolds with singularities and boundaries.
- What about infinite-dimensional data? Cause images come from infinite-dimensional continuous functions.

[Emile says: would be nice to elaborate on the above, and to point to specific instances where assuming it has some "catastrophic" consequences. that could go into a motivating examples section]

1.1 Related Works

Energy-based models Song and Kingma [2021], Du and Mordatch [2019], LeCun et al. [2006], Zhu et al. [2024], Du et al. [2021, 2020]

Compositional generation Su et al. [2024], Liu et al. [2021, 2022, 2023], Du and Kaelbling [2024], Bagautdinov et al. [2018]

Occam's razor Elmoznino et al. [2025], Rasmussen and Ghahramani [2000]

Bakry-Emery

Dirichlet Form

Example: Vision Data Each 2d image (or 3d scene or a point cloud) can be modeled $f : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ as a function from coordinates to RGB (or XYZ) values. That's how NeRFS work:

Neural Radiance Fields (NeRFs). Mildenhall et al. [2021], Pumarola et al. [2021], Xu et al. [2021], Wang et al. [2022], Rabby and Zhang [2024], Reiser et al. [2021], Müller et al. [2022], Tancik et al. [2021], Yu et al. [2021], Chen et al. [2022] [Emile says: this actually looks very relevant. would be worth doing a deeper dive into this] A *Neural Radiance Field* models a 3D scene as a continuous function

$$f_{\theta} : (\mathbf{x}, \mathbf{d}) \mapsto (\sigma, \mathbf{c}),$$

where

- $\mathbf{x} \in \mathbb{R}^3$ denotes a spatial location,
- $\mathbf{d} \in \mathbb{S}^2$ denotes a viewing direction,
- $\sigma \in \mathbb{R}_{\geq 0}$ is the volume density at $\mathbf{x}$,
- $\mathbf{c} \in \mathbb{R}^3$ is the emitted RGB color in direction $\mathbf{d}$.

The key idea is that the scene is represented neither as a mesh, nor as a voxel grid, nor as a point cloud, but as a *continuous radiance field* parameterized by a neural network.

In a similar way, each 2d image can be modeled as a function from coordinates to RGB values: $f : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ and the embedding of this function would be the weights of the corresponding ReLU network. The regular manifold hypothesis doesn't properly capture this structure that:

- In a union of manifolds
- Has some hierarchical structure.
- Has permutation-invariance with respect to weights in each layer.

Also, importantly, each image f is not a smooth function: its a combination of a "simple" in some sense partition of the image into smooth enough regions.

69 2 Set Up

70 We work in an abstract setting that captures the common statistical belief that observed data
 71 $x_1, \dots, x_n \in X$ concentrate on a low-dimensional structure, even when the ambient space X
 72 is high-dimensional. Concretely, we fix a *reference* distribution ν_0 on X (e.g. Gaussian or uniform)
 73 encoding our prior knowledge about X , and posit that the data are drawn from a *target* distribution μ
 74 that is far more concentrated — e.g. supported on a d -dimensional manifold $M \subset \mathbb{R}^D$ with $d \ll D$.
 75 The benefit is quantitative: greater concentration of μ yields faster learning rates for functions
 76 $f : X \rightarrow \mathbb{R}$.

77 **Differential structure.** To differentiate functions on X we fix an *algebra of test functions* $\mathcal{A}$ (a
 78 subset of $\mathbb{R}^X$ closed under pointwise products) and a *derivation*

$$D : \mathcal{A} \longrightarrow \mathcal{M}$$

79 into a module $\mathcal{M}$ equipped with a pointwise inner product $\langle \cdot, \cdot \rangle_x$, satisfying the Leibniz rule $D(fg) =$
 80 $fDg + gDf$. This gives rise to the *carré du champ* operator

$$\Gamma(f, g)(x) := \langle Df(x), Dg(x) \rangle_x,$$

81 which generalizes the square-gradient $|\nabla f|^2$ to the abstract setting.

82 **Dirichlet form and Laplacian.** Given any σ -finite measure ν on $(X, \mathcal{F})$, the carré du champ
 83 integrates to a *Dirichlet form*:

84 **Definition 2.1** (Symmetric Dirichlet Space [Bakry et al., 2014, §3.4]). Let $(X, \mathcal{F}, \mu)$ be a σ -finite
 85 measure space. A *Dirichlet form* on $L^2(\mu)$ is a pair $(\mathcal{E}, \mathcal{D}(\mathcal{E}))$ where:

- 86 1. $\mathcal{D}(\mathcal{E}) \subset L^2(\mu)$ is a dense linear subspace;
- 87 2. $\mathcal{E} : \mathcal{D}(\mathcal{E}) \times \mathcal{D}(\mathcal{E}) \rightarrow \mathbb{R}$ is a symmetric bilinear form;
- 88 3. **(Non-negativity)** $\mathcal{E}(f, f) \geq 0$ for all $f \in \mathcal{D}(\mathcal{E})$;
- 89 4. **(Closedness)** $\mathcal{D}(\mathcal{E})$ is complete under $\|f\|_{\mathcal{E}}^2 := \|f\|_{L^2(\mu)}^2 + \mathcal{E}(f, f)$;
- 90 5. **(Markov property)** for every $f \in \mathcal{D}(\mathcal{E})$, the truncation $\tilde{f} := (0 \vee f) \wedge 1$ satisfies $\tilde{f} \in \mathcal{D}(\mathcal{E})$
 91 and $\mathcal{E}(\tilde{f}, \tilde{f}) \leq \mathcal{E}(f, f)$.

92 The quadruple $(X, \mathcal{F}, \mu, \mathcal{E})$ is called a *Symmetric Dirichlet Space*.

93 In our setting the Dirichlet form is

$$\mathcal{E}_\nu(f, g) = \int_X \Gamma(f, g) d\nu,$$

94 so $\mathcal{E}_\nu(f, f)$ is the total gradient energy of f under ν . The Markov property is automatic from the
 95 Leibniz rule and the truncation behaviour of Γ .

96 **Definition 2.2** (Generalized Laplacian). To any Dirichlet form $(\mathcal{E}_\nu, \mathcal{D}(\mathcal{E}_\nu))$ on $L^2(\nu)$ there corre-
 97 sponds a unique self-adjoint operator L_ν on $L^2(\nu)$ satisfying

$$\mathcal{E}_\nu(f, g) = \langle f, L_\nu g \rangle_{L^2(\nu)}, \quad \forall f \in \mathcal{D}(\mathcal{E}_\nu), g \in \mathcal{D}(L_\nu).$$

98 L_ν admits a spectral decomposition $L_\nu \phi_j = \lambda_j \phi_j$ with

$$0 = \lambda_0 \leq \lambda_1 \leq \lambda_2 \leq \dots, \quad \langle \phi_j, \phi_k \rangle_{L^2(\nu)} = \delta_{jk}.$$

99 The Markov property ensures that e^{-tL_ν} is a Markov semigroup: starting from any initial distribution
 100 ν_0 on X , the heat flow $e^{-tL_\nu} \nu_0$ converges to ν as $t \rightarrow \infty$, confirming that ν is the invariant measure
 101 of the diffusion. The *iterated carré du champ* Γ_2^ν measures how Γ evolves under this semigroup:

$$\Gamma_2^\nu(f, g) = \frac{1}{2}(L_\nu \Gamma(f, g) - \Gamma(f, L_\nu g) - \Gamma(g, L_\nu f)), \quad \Gamma_2^\nu(f) := \Gamma_2^\nu(f, f).$$

102 The Bakry–Émery condition $\Gamma_2^\nu(f) \geq \rho \Gamma(f)$ ($\rho \in \mathbb{R}$) encodes a lower bound on the Ricci curvature
 103 of the underlying space.

104 **Spectral dimension.** The eigenvalue asymptotics of L_ν encode the intrinsic dimension of the
 105 learning space, providing the key rate parameter throughout the paper.

106 **Definition 2.3** (Effective Dimension). The *effective dimension* of $(X, \mathcal{F}, \nu, \mathcal{E}_\nu)$ is the quantity
 107 $d_{\text{eff}} \in (0, \infty)$ characterised by any of the following equivalent asymptotics:

$$\begin{aligned} \lambda_j &\asymp j^{2/d_{\text{eff}}} && \text{(eigenvalue growth),} \\ N(\lambda) := \#\{j : \lambda_j \leq \lambda\} &\asymp \lambda^{d_{\text{eff}}/2} && \text{(counting function),} \\ \text{tr } e^{-tL_\nu} &\sim t^{-d_{\text{eff}}/2} && \text{(heat-kernel trace, as } t \rightarrow 0^+). \end{aligned}$$

108 For a compact d -dimensional Riemannian manifold M , Weyl's law gives $d_{\text{eff}} = d$, confirming that
 109 the effective dimension recovers the geometric dimension in the classical setting.

110 3 Effective Dimension Hypothesis

111 **Assumption 3.1** (Effective Dimension Hypothesis). All the observed data $x_1, \dots, x_n \in X$ come
 112 from a learning space $(X, \mathcal{A}, \mu)$ with small effective dimension even if the ambient space $(X, \mathcal{A}, \nu_0)$
 113 has large effective dimension where ν_0 is the base distribution on X :

$$d_{\text{eff}}(\mu) \ll d_{\text{eff}}(\nu_0)$$

114 4 Relevant Things

Definition 4.1 (Curvature–Dimension Condition, Ledoux [2000] (Definition 1.2)). For $\rho \in \mathbb{R}, N \in (0, \infty)$, L_ν satisfies $\text{CD}(\rho, N)$ if for all $f \in \mathcal{A}$

$$\Gamma_2^\nu(f, f) \geq \rho \Gamma(f, f) + \frac{1}{N} (L_\nu f)^2$$

115 Intuitively: $\rho \sim$ curvature parameter, $N \sim$ dimension

116 **Example 4.2.** Examples of generators L_ν that satisfy the CD condition:

- 117 • An n -dimensional complete Riemannian manifold (M, g) with Ricci curvature bounded
- 118 below, or rather the Laplacian Δ on M , satisfies the inequality $\text{CD}(R, n)$ with R the
- 119 infimum of the Ricci tensor and n the geometric dimension.
- 120 • The classical Laplacian Δ on $\mathbb{R}^n$ satisfies $\text{CD}(0, n)$,
- 121 • the Laplace–Beltrami operator on the unit sphere S^n is of curvature dimension $\text{CD}(n-1, n)$
- 122 since curvature is constant and equal to $n-1$ in this case.

123 4.1 Sub-Exponential Concentration

Definition 4.3 (Log-Sobolev Inequality). Given μ for there is a constant C_{LS} such that the following holds for any $f : X \rightarrow \mathbb{R}$

$$\text{Ent}_\mu(f^2) \leq 2C_{LS} \int \|Df\|^2 d\mu$$

124 where $\text{Ent}_\mu(g) = \int g \log g d\mu - \int g d\mu \cdot \log \int g d\mu$ for $g \geq 0$.

125 **Lemma 4.4.** Assuming $\Gamma_2^\mu(f, f) \geq \rho \Gamma(f, f)$ for all $f \in \mathcal{A}$ then $C_{LS} \leq \frac{1}{\rho}$

126 Consequences:

Theorem 4.5 (Hypercontractivity, Gross 1975). The LSI is equivalent to hypercontractivity of the semigroup: for $1 < p < q < \infty$,

$$\|P_t f\|_{L^q(\mu)} \leq \|f\|_{L^p(\mu)} \quad \text{when } t \geq \frac{C_{LS}}{2} \log \frac{q-1}{p-1}$$

Theorem 4.6 (Sub-Gaussian concentration). For f with $\Gamma(f, f) \leq 1, \mu$ -a.e., the Herbst argument gives

$$\mu(f - \mathbb{E}_\mu f \geq t) \leq \exp\left(-\frac{t^2}{2C_{LS}}\right)$$

127 The smaller C_{LS} – the better.

128 4.2 Effective Dimension

129 **Corollary 4.7.** If ν satisfies $\text{CD}(\rho, N)$ then $d_{\text{eff}} \leq N$

130 5 Supervised Learning: Learning a Function From Data

131 Adding a prior $p(f) \propto e^{-E(f)}$ over functions $f : X \rightarrow \mathbb{R}$ the supervised learning problem is
 132 equivalent to the following optimization problem:

$$\hat{f}_n = \min_f \sum_{i=1}^n \|f(x_i) - y_i\|^2 + E(f)$$

133 Note, that $E(f)$ defines sublevel sets $\mathcal{F}_r = \{f : E(f) \leq r\}$. The generalization error of $\hat{f}_n$
 134 decomposes as:

$$\mathbb{E} \left\| f^* - \hat{f}_n \right\|_{L^2}^2 \leq \underbrace{\inf_{f: E(f) \leq r} \|f - f^*\|^2}_{\text{bias: how well } \mathcal{F}_r \text{ approximates } f^*} + \underbrace{\varepsilon_n^2(r)}_{\text{variance: estimation error within } \mathcal{F}_r}$$

Definition 5.1 (Sublevel set). Given a convex functional $E : \mathcal{A} \rightarrow \mathbb{R} \cup \{+\infty\}$, a radius $r > 0$, and a metric d on $\mathcal{A}$, the sublevel set is:

$$\mathcal{B}_r(E) = \{f \in \mathcal{A} : E(f) \leq r\}$$

135 **Definition 5.2** (Covering number). The ϵ -covering number is:

$$N(\epsilon, \mathcal{B}_r(E), d) = \min \left\{ m \in \mathbb{N} : \exists f_1, \dots, f_m \in \mathcal{A} \text{ s.t. } \mathcal{B}_r(E) \subseteq \bigcup_{i=1}^m \{g : d(g, f_i) \leq \epsilon\} \right\}$$

136 The natural metric in our setting is $d(f, g) = \|f - g\|_{L^2(\mu)} = \left(\int_X (f - g)^2 d\mu \right)^{1/2}$

Definition 5.3 (Metric entropy).

$$\mathcal{H}(\epsilon, \mathcal{B}_r(E), d) = \log N(\epsilon, \mathcal{B}_r(E), d)$$

137 The smaller the metric entropy – the less samples we need to learn a function from data.

138 Lets consider 3 simple enough cases of what $E(f)$ can be:

139 1. Given convex $g : \mathbb{R} \rightarrow \mathbb{R}$ define $E(f) = \int_X g(f(x)) d\mu(x)$

140 2. $E(f) = \mathcal{E}_\mu(f, f)$

141 3. Given $\phi : \mathbb{R} \rightarrow \mathbb{R}$ define $E(f) = \langle f, \phi(L_\mu) f \rangle$

142 **Proposition 5.4.** For $E(f) = \int_X g(f(x)) d\mu(x)$ The covering number of $\mathcal{B}_r(E)$ is at most

143 $N(r, E) = \frac{1}{r} \int_0^r \mathcal{H}(\epsilon, \mathcal{B}_r(E), d) d\epsilon.$

144 As we introduced Definition 4.1 that captures the "compactness" of the measure, we can say now that
 145 $\mu \ll \nu$ if

$$\begin{aligned} \mu &\in CD(\rho_\mu, N_\mu), \nu \in CD(\rho_\nu, N_\nu) \\ &\text{and} \\ \rho_\mu &> \rho_\nu, N_\mu < N_\nu \end{aligned}$$

146 The following theorem captures what density $\frac{d\mu}{d\nu}$ "simplifies" the measure.

147 **Theorem 5.5** (Perturbation of $CD(\rho, N)$). *[TODO: double check, add a proof, better notation]*

148 Let $(X, \mathcal{A}, \Gamma, \mu_0)$ be a Markov diffusion structure with generator L_0 satisfying $CD(\rho, N)$:

$$\Gamma_2^{(0)}(f, f) \geq \rho \Gamma(f, f) + \frac{1}{N} (L_0 f)^2 \quad \text{for all } f \in \mathcal{A}$$

149 Let $U \in \mathcal{A}$ satisfy:

150 • *Abstract κ -convexity:* $\Gamma(f, \Gamma(U, f)) - \frac{1}{2}\Gamma(U, \Gamma(f, f)) \geq \kappa\Gamma(f, f)$ for all f

151 • *Bounded abstract gradient:* $\Gamma(U, U) \leq M$ μ_0 -a.e.

152 Let $d\mu = Z^{-1}e^{-U}d\mu_0$ with $L_U f = L_0 f - \Gamma(U, f)$. Then for any $\varepsilon \in (0, 1)$, (L_U, Γ, μ) satisfies
153 CD (ρ', N') with

$$\rho' = \rho + \kappa - \frac{(1/\varepsilon - 1)M}{N}, \quad N' = \frac{N}{1 - \varepsilon}$$

154 In particular, $C_{LS}(\mu) \leq \frac{1}{\rho'}$ whenever $\rho' > 0$.

6 Diffusion Models

[TODO: Add noise schedule maybe.]

Forward process: $\mu \rightarrow \mu_0$ Density $p_t = \mu_t / \mu_0$ evolves according to the diffusion equation:

$$\frac{\partial p_t}{\partial t} = L_{\mu_0} p_t \implies p_t = e^{tL_{\mu_0}} \mu_0$$

Backward process: $\mu_0 \rightarrow \mu$ Density p_t evolves according to the diffusion+drift equation:

$$\frac{\partial p_t}{\partial t} = L_{\mu_0} p_t - \nabla \cdot (p_t \nabla \log p_t)$$

To be able to reverse it back, we need to learn $\nabla \log p_t$. Let us look at the very first δ step of the forward process (corresponding to the very last δ step of the backward process) and learn a score function for that step $s(\cdot) = \nabla \log p_T = \nabla \log \frac{d\mu}{d\mu_0}$. The minimization problem is:

$$\min_s \underbrace{\mathbb{E}_{p_T} [\|s(\cdot) - \nabla \log p_T\|^2]}_{\text{data term}} + \underbrace{R(s)}_{\text{regularizer term}} \quad (1)$$

While the data term is straightforward, the regularizer term is to be chosen. Assuming $d\mu = Z^{-1}e^{-U}d\mu_0 \Rightarrow \nabla \log p_T = -DU$. Our prior (hope) for U is that it concentrates the measure better, specifically, that by Definition 4.1 $U \in \mathcal{A}$ satisfies:

- κ -convexity: $\Gamma(f, \Gamma(U, f)) - \frac{1}{2}\Gamma(U, \Gamma(f, f)) \geq \kappa\Gamma(f, f)$ for all f
- Bounded gradient: $\Gamma(U, U) \leq M$ μ_0 -a.e.

Define the abstract Hessian as a map from 1-forms to bilinear forms on $D(\mathcal{A})$:

$$H_v(f) := \langle Df, D\langle v, Df \rangle \rangle - \frac{1}{2}\langle v, D\langle Df, Df \rangle \rangle$$

This is linear in v : $H_{\alpha v_1 + \beta v_2}(f) = \alpha H_{v_1}(f) + \beta H_{v_2}(f)$. Note, that the term $\langle v, v \rangle \leq M$ is quadratic in v .

Hence, the following set is convex:

$$\mathcal{C}_{\kappa, M} = \{v \in D(\mathcal{A}) : H_v(f) \geq \kappa\Gamma(f, f) \text{ and } \langle v, v \rangle \leq M \quad \forall f \in \mathcal{A}\}$$

First let us compute its covering number.

Lemma 6.1. The covering number of $\mathcal{C}_{\kappa, M}$ is at most $N(M, \kappa) = \frac{M}{\kappa}$.

Lemma 6.2 (Very rough sketch). Define $W_{\kappa, M}$ as [TODO: a fixed set of very cool points from $\mathcal{C}_{\kappa, M}$]. Then for every $v \in \mathcal{C}_{\kappa, M}$ the following problem has a unique solution:

$$\min_{\omega \in \mathfrak{M}_+(W_{\kappa, M})} \|\omega\|_{TV} \text{ such that } \int_{w \in W_{\kappa, M}} w d\omega(w) = f$$

Hence, if we know the set $W_{\kappa, M}$ then Equation (1) can be rewritten as:

$$\min_{\omega \in \mathfrak{M}_+(W)} \mathbb{E}_{p_T} \left[\left\| \int_{w \in W_{\kappa, M}} w d\omega(w) - \nabla \log p_T \right\|^2 \right] + \|\omega\|_{TV}$$

6.1 References

<https://djalil.chafai.net/blog/2023/01/12/log-sobolev-and-bakry-emery/>

176 **7 Energy-Based Model Perspective on Effective Dimension**

177 Bakry-Emery theorem??

178 need framework that shows that pure noise, fractals, hierarchical representations should be hard to
179 learn

180 **8 Experiments**

181 **A learning curve experiment:** estimate a function class regularized by the energy, and show that
182 lower finite-scale complexity predicts sample complexity better than ambient dimension.

183

184 **A 2D synthetic suite** one smooth manifold, a union of two manifolds with intersection, and a
185 stratified example with a singular point. Show that asymptotic Weyl dimension is not the right
186 discriminator, while $d_\mu(t)$ at finite t separates them.

187 **9 Conclusion**

References

- Timur Bagautdinov, Chenglei Wu, Jason Saragih, Pascal Fua, and Yaser Sheikh. Modeling facial geometry using compositional VAEs. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3877–3886, 2018.
- Dominique Bakry, Ivan Gentil, and Michel Ledoux. *Analysis and Geometry of Markov Diffusion Operators*, volume 348 of *Grundlehren Der Mathematischen Wissenschaften*. Springer International Publishing, Cham, 2014. ISBN 978-3-319-00226-2 978-3-319-00227-9. doi: 10.1007/978-3-319-00227-9.
- Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. Tensorf: Tensorial radiance fields. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXII*, page 333–350, Berlin, Heidelberg, 2022. Springer-Verlag. ISBN 978-3-031-19823-6. doi: 10.1007/978-3-031-19824-3_20. URL https://doi.org/10.1007/978-3-031-19824-3_20.
- Yilun Du and Leslie Kaelbling. Position: compositional generative modeling: a single model is not all you need. In *Proceedings of the 41st International Conference on Machine Learning, ICML’24*. JMLR.org, 2024.
- Yilun Du and Igor Mordatch. Implicit generation and modeling with energy based models. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Yilun Du, Shuang Li, and Igor Mordatch. Compositional visual generation with energy based models. *Advances in Neural Information Processing Systems*, 33:6637–6647, 2020.
- Yilun Du, Shuang Li, Yash Sharma, Joshua B. Tenenbaum, and Igor Mordatch. Unsupervised learning of compositional energy concepts. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021.
- Eric Elmoznino, Tom Marty, Tejas Kasetty, Leo Gagnon, Sarthak Mittal, Mahan Fathi, Dhanya Sridhar, and Guillaume Lajoie. In-context learning and occam’s razor, 2025. URL <https://arxiv.org/abs/2410.14086>.
- Yann LeCun, Sumit Chopra, Raia Hadsell, M Ranzato, Fugie Huang, et al. A tutorial on energy-based learning. *Predicting structured data*, 1(0), 2006.
- Michel Ledoux. The geometry of Markov diffusion generators. *Annales de la Faculté des sciences de Toulouse : Mathématiques*, 9(2):305–366, 2000. ISSN 2258-7519.
- Nan Liu, Shuang Li, Yilun Du, Joshua B. Tenenbaum, and Antonio Torralba. Learning to compose visual relations. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021.
- Nan Liu, Shuang Li, Yilun Du, Antonio Torralba, and Joshua B. Tenenbaum. Compositional visual generation with composable diffusion models. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XVII*, page 423–439, Berlin, Heidelberg, 2022. Springer-Verlag. ISBN 978-3-031-19789-5. doi: 10.1007/978-3-031-19790-1_26. URL https://doi.org/10.1007/978-3-031-19790-1_26.
- Nan Liu, Yilun Du, Shuang Li, Joshua B. Tenenbaum, and Antonio Torralba. Unsupervised compositional concepts discovery with text-to-image generative models. In *ICCV*, pages 2085–2095, 2023.
- Ben Mildenhall, Pratul P Srinivasan, Matthew Tancik, Jonathan T Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. *Communications of the ACM*, 65(1):99–106, 2021.
- Thomas Müller, Alex Evans, Christoph Schied, and Alexander Keller. Instant neural graphics primitives with a multiresolution hash encoding. *ACM Trans. Graph.*, 41(4), July 2022. ISSN 0730-0301. doi: 10.1145/3528223.3530127. URL <https://doi.org/10.1145/3528223.3530127>.

235 Albert Pumarola, Enric Corona, Gerard Pons-Moll, and Francesc Moreno-Noguer. D-nerf: Neural
236 radiance fields for dynamic scenes. In *Proceedings of the IEEE/CVF conference on computer
237 vision and pattern recognition*, pages 10318–10327, 2021.

238 AKM Shahariar Azad Rabby and Chengcui Zhang. Beyondpixels: A comprehensive review of the
239 evolution of neural radiance fields, 2024. URL <https://arxiv.org/abs/2306.03000>.

240 Carl Rasmussen and Zoubin Ghahramani. Occam’s razor. *Advances in neural information processing
241 systems*, 13, 2000.

242 Christian Reiser, Songyou Peng, Yiyi Liao, and Andreas Geiger. Kilonerf: Speeding up neural
243 radiance fields with thousands of tiny mlps. In *2021 IEEE/CVF International Conference on
244 Computer Vision (ICCV)*, pages 14315–14325, 2021. doi: 10.1109/ICCV48922.2021.01407.

245 Yang Song and Diederik P. Kingma. How to train your energy-based models, 2021.

246 Jocelin Su, Nan Liu, Yanbo Wang, Joshua B. Tenenbaum, and Yilun Du. Compositional image
247 decomposition with diffusion models. In *Proceedings of the 41st International Conference on
248 Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 46823–46842.
249 PMLR, 21–27 Jul 2024.

250 Matthew Tancik, Ben Mildenhall, Terrance Wang, Divi Schmidt, Pratul P. Srinivasan, Jonathan T.
251 Barron, and Ren Ng. Learned initializations for optimizing coordinate-based neural representations.
252 In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2845–
253 2854, 2021. doi: 10.1109/CVPR46437.2021.00287.

254 Zirui Wang, Shangzhe Wu, Weidi Xie, Min Chen, and Victor Adrian Prisacariu. Nerf–: Neural
255 radiance fields without known camera parameters, 2022. URL [https://arxiv.org/abs/2102.
256 07064](https://arxiv.org/abs/2102.07064).

257 Hongyi Xu, Thiemo Alldieck, and Cristian Sminchisescu. H-nerf: Neural radiance fields for rendering
258 and temporal reconstruction of humans in motion. In *Advances in Neural Information Processing
259 Systems (NeurIPS)*, 2021. URL <https://arxiv.org/pdf/2110.13746.pdf>.

260 Alex Yu, Vickie Ye, Matthew Tancik, and Angjoo Kanazawa. pixelnerf: Neural radiance fields from
261 one or few images. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition
262 (CVPR)*, pages 4576–4585, 2021. doi: 10.1109/CVPR46437.2021.00455.

263 Yaxuan Zhu, Jianwen Xie, Ying Nian Wu, and Ruiqi Gao. Learning energy-based models by
264 cooperative diffusion recovery likelihood. In *The Twelfth International Conference on Learning
265 Representations*, 2024.

267 B Recovering the Manifold Hypothesis

268 We show that the manifold hypothesis is recovered after coarsening the effective dimension hypothesis
269 in Theorem B.1.

270 **Theorem B.1.** *Let $M \subset \mathbb{R}^D$ be a compact embedded C^2 submanifold with tubular radius r_0 , with
271 the nearest-point projection $\pi : T_{r_0}(M) \rightarrow M$. Fix $0 < \epsilon < r_0$. Let μ be a Borel probability
272 measure supported on the tangent space $T_\epsilon(M) := \{x \in \mathbb{R}^D : \text{dist}(x, M)\} \leq \epsilon\}$, and define the manifold-
273 supported measure $\bar{\mu} := \pi_{\#}\mu$. Then, for every $p \geq 1$, we have $W_p(\mu, \bar{\mu}) \leq \epsilon$ and for every
274 L -Lipschitz observable function $f : \mathbb{R}^D \rightarrow \mathbb{R}$, we have that*

$$\left| \int f d\mu - \int f d\bar{\mu} \right| \leq L\epsilon.$$

275 *Proof.* Let $X \sim \mu$ and set $Y := \pi(X)$. Since $\bar{\mu} := \pi_{\#}\mu$, we have $Y \sim \bar{\mu}$. Since $X \in T_\epsilon(M)$ almost
276 surely, and $\pi(X)$ is the nearest point on M , we have $\|X - Y\| = \|X - \pi(X)\| = \text{dist}(X, M) \leq \epsilon$,
277 almost surely. Next, we consider the coupling $\gamma = (\text{Id}, \pi)_{\#}\mu \in \Pi(\mu, \bar{\mu})$ of μ and $\bar{\mu}$, and γ is
278 supported on pairs x, y with $\|x - y\| \leq \epsilon$. Therefore, for $1 \leq p < \infty$, we have

$$\begin{aligned} W_p(\mu, \bar{\mu})^p &\leq \int \|x - y\|^p d\gamma(x, y) \\ &= \int \|x - \pi(x)\|^p d\mu(x) \leq \epsilon^p, \end{aligned}$$

279 which implies $W_p(\mu, \bar{\mu}) \leq \epsilon$. Moreover, we have

$$W_\infty(\mu, \bar{\mu}) \leq \text{ess sup}_{x \sim \mu} \|x - \pi(x)\| \leq \epsilon.$$

280 Next, since $\bar{\mu} = \pi_{\#}\mu$, we have $\int f d\bar{\mu} = \int f(\pi(x)) d\mu(x)$. Therefore, we have

$$\begin{aligned} \left| \int f d\mu - \int f d\bar{\mu} \right| &= \left| \int (f(x) - f(\pi(x))) d\mu(x) \right| \\ &\leq \int \|f(x) - f(\pi(x))\| d\mu(x) \\ &\leq \int L\|x - \pi(x)\| d\mu(x) \\ &\leq L\epsilon \int d\mu(x) := L\epsilon, \end{aligned}$$

281 which proves the theorem. □

282 **Corollary B.2.** *Let $F_\delta := \{f : \text{Lip}(f) \leq 1/\delta\}$. By Theorem B.1, we have that $\sup_{f \in F_\delta} \left| \int f d\mu - \int f d\bar{\mu} \right| \leq \epsilon/\delta$. Therefore, once $\delta \gg \epsilon$, the tube supported distribution μ and the manifold-supported
283 distribution $\bar{\mu}$ are indistinguishable by coarse observations. Therefore, the manifold hypothesis is
284 recovered upon coarsening the effective dimension hypothesis.*

286 **Remark B.3.** The proof of Theorem B.1 is based on the metric space, not on any of its spectral
287 properties. Moreover, the manifold structure is only used to guarantee the existence of the projection
288 π , and the same argument would also follow for any measurable set $S \subset \mathbb{R}^D$ equipped with a
289 measurable map $\pi : T_\epsilon(S) \rightarrow S$ such that $\|x - \pi(x)\| \leq \epsilon$, which allows the result to extend beyond
290 smooth manifolds. Finally, the Lipschitz bound in Theorem B.1 is tight in general. For this, take
291 $M = \{0\} \subset \mathbb{R}^D$, $\mu = \delta_\epsilon$, $\bar{\mu} = \delta_0$, and $f(x) = Lx$. Then, $\left| \int f d\mu - \int f d\bar{\mu} \right| = L\epsilon$.